



# Process

Distributed Systems IT332



# Outline

- Threads
- Virtualization
- Server
- Code Migration

# Processes and Threads

## Process

- An instance of a computer program that is being executed on one of the operating system's (virtual) processors .
  - While a program itself is just a passive collection of instructions, a process is the actual execution of those instructions
- A process has a virtual address space, executable code, open handles to system objects, a security context, a unique identifier, environment variables, a priority class, and at least one thread of execution.

# Processes and Threads

## Threads

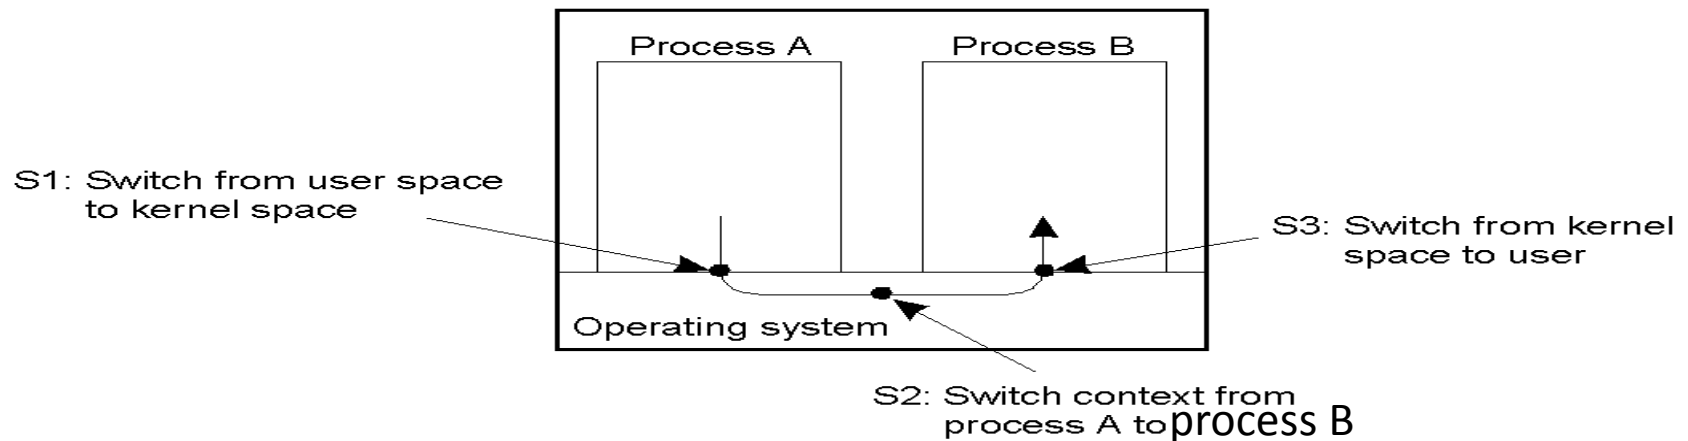
- Each process is started with a single thread, often called the primary thread, but can create additional threads from any of its threads.
- Thread (lightweight process)
  - A thread is a path of execution within a process. A thread executes its own piece of code, independently from other threads.
  - A way for a process to split itself into two or more simultaneously running tasks.
  - A thread maintains only the minimum information to allow a CPU to be shared by several threads. ?
  - In particular, a thread context often consists of nothing more than the CPU context.

# Processes and Threads

## Advantages of Threads

- Allow for parallel computation.
- Switching between threads is faster than switching between process.
- Creating and destroying threads is cheaper than creating and destroying a process since that:
  - Threads use very little resources of an operating system in which they are working.
  - Threads do not need new address space, global data, program code or operating system resources and can share common data
- Easier to structure many applications as a collection of cooperating threads
- Higher performance compared to multiple processes since switching between threads takes less time.

# Thread Usage in Non-Distributed Systems



Context switching as the result of IPC.

- The major drawback of all IPC mechanisms is that communication often requires extensive context switching, shown at three different points.

# Thread Usage in Non-Distributed Systems

- Example: A spreadsheet program where the different cells are dependent, when a user changes the value in a single cell, such a modification can trigger a large series of computations:
  - If there is only a single thread of control, computation cannot proceed while the program is waiting for input. Likewise, it is not easy to provide input while dependencies are being calculated.
- → The easy solution is to have at least two threads of control: one for handling interaction with the user and one for updating the spreadsheet. In the mean time, a third thread could be used for backing up the spreadsheet to disk while the other two are doing their work.

# Threads in Distributed Systems

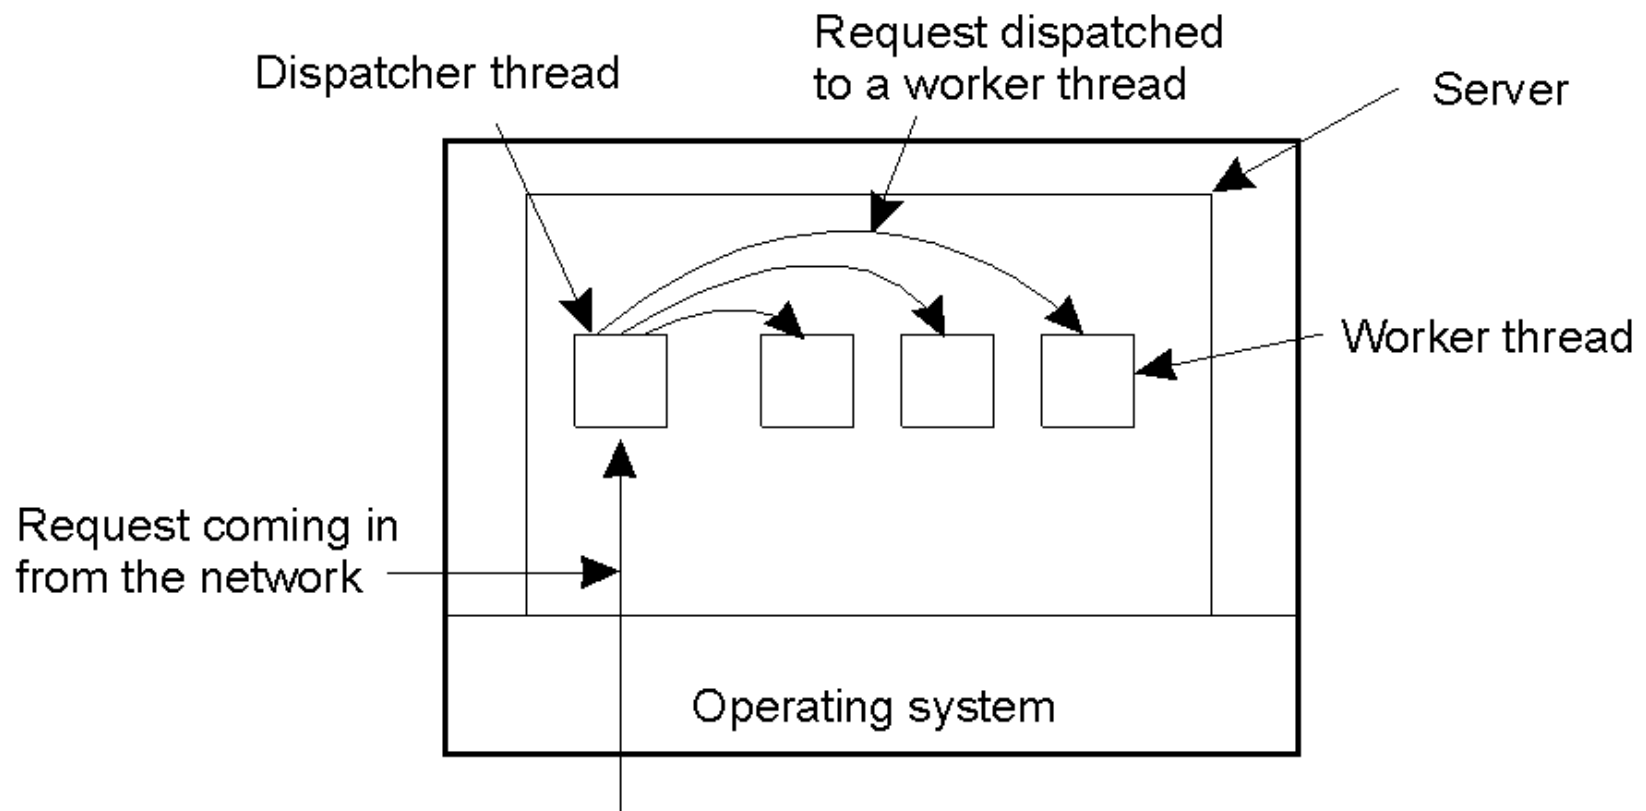
- Used to express communication in the form of multiple logical connections at the same time.
- An important property of threads is that they can provide a convenient means of allowing blocking system calls without blocking the entire process in which the thread is running.
- A main contribution of threads in distributed systems is that they allow clients and servers to be constructed such that communication and local processing can overlap, resulting in a high level of performance.
- Attractive to use in distributed systems
  - Multithreaded client
  - Multithreaded server



# Multithreaded Clients

- Can be used to hide delays/latencies in network communications, by initiating communication and immediately proceeding with something else.
- Example: web browsers such as IE are multi-threaded
  - A web browser can start up several threads: once the main HTML file has been fetched, separate threads can be activated to take care of fetching the other parts. Each thread sets up a separate connection to the server and pulls in the data. One for downloading the HTML source of the page, one each for images on the page, one each for animations/applets etc.
- Replicated web servers along with multi-threaded clients can result in shorter download times.

# Multithreaded Servers



A multithreaded server organized in a dispatcher/worker model

# Multithreaded Servers


Model	Characteristics
Threads	Parallelism, blocking system calls
Single-threaded process	No parallelism, blocking system calls
Finite-state machine	Parallelism, nonblocking system calls

## Three ways to construct a server

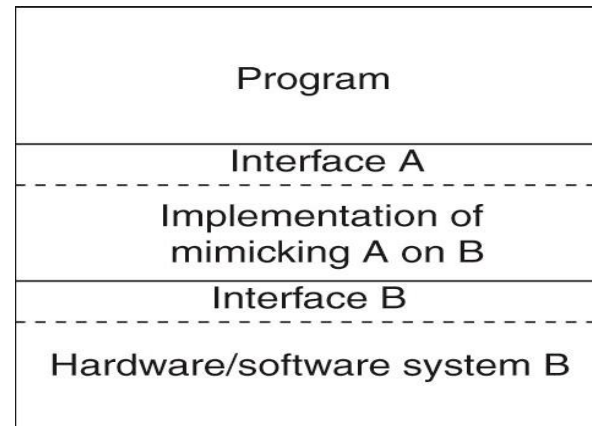
- Blocking system calls make programming easier and parallelism improves performance.
- The single-threaded server retains the ease and simplicity of blocking system calls, but gives up performance.
- The finite-state machine approach achieves high performance through parallelism, and uses nonblocking calls, thus is hard to program.

# Virtualization

- Virtualization is the creation of a virtual (rather than actual) version of something, such as a hardware platform, resources, operating system, a storage device or network resources.
- Introduced in 1970s. IBM have applied this technique very successfully for a long time on the IBM 370 mainframes (and their successors): to offer a virtual machine to which different OS had been ported.



(a) General organization between a program, interface, and system.




(b) General organization of virtualizing system A on top of system B.

# Role of Virtualization

- Helps with scalability and better utilization of hardware resources.
- Allows legacy software to run on expensive mainframe hardware.
  - Suppose you want to run an existing application on a new computer
- Runs multiple different operating systems at the same time.
- Provides a high degree of portability and flexibility.

# Architectures of Virtual Machines



# Architectures of Virtual Machines

## ➤ Four types of interfaces:

- An interface between the hardware and software, consisting of machine instructions
  - Can be invoked by any program.
- An interface between the hardware and software, consisting of machine instructions
  - can be invoked only by privileged programs, such as an operating system.
- An interface consisting of system calls as offered by an operating system.
- An interface consisting of library calls,
  - known as an application programming interface (API).

# Architectures of Virtual Machines

➤ Virtualization can be implemented at two levels.

➤ **Process Virtual Machine:**

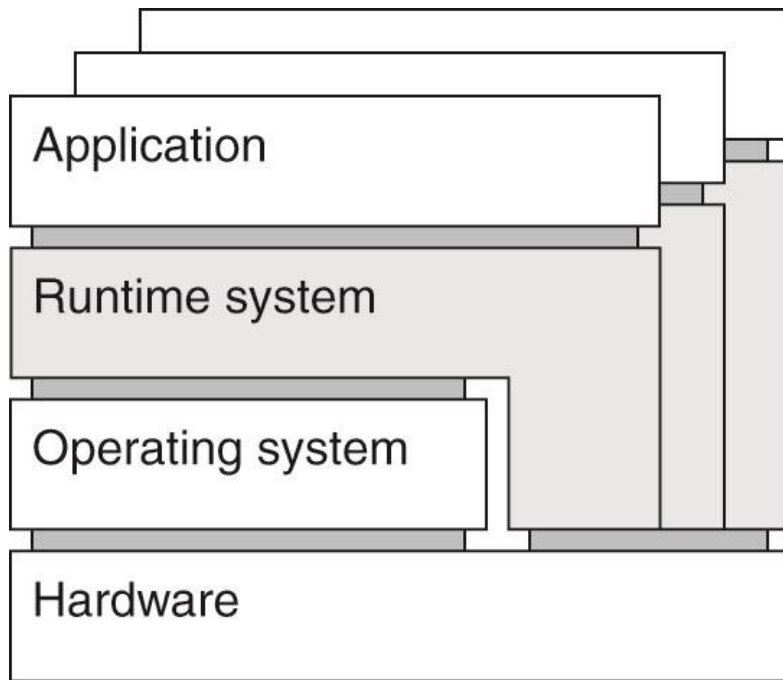
- Virtualization is done essentially only for executing a single process (program).
- An abstract instruction set that is to be used for executing applications.
- For example: Java runtime, Windows emulation (Wine) on Unix/Linux/MacOS.

➤ **Virtual Machine Monitor:**


- Composed of the host OS and the virtualization software.
- A layer completely shielding the original hardware but offering the complete instruction set of that same (or other hardware) as an interface.
- Provides a further decoupling between hardware and software allowing moving complete environment from one machine to another.
- Makes it possible to have multiple instances of different operating systems run simultaneously and concurrently on the same platform.
- Examples: Vmware, VirtualBox, Xen, VirtualPC, Parallels etc.



# Architectures of Virtual Machines



(a)



(b)

A process virtual machine, with multiple instances of (application, runtime) combinations. (a)

(b) A virtual machine monitor, with multiple instances of (applications, OS) combinations.

# Servers

- There are several ways to organize servers:
  - In the case of an iterative server, the server itself handles the request and, if necessary, returns a response to the requesting client.
  - A concurrent server can be multi-threaded or multi-process. It does not handle the request itself, but passes it to a separate thread or another process, after which it immediately waits for the next incoming request.
- A server can be **stateless** or **stateful**.
  - A stateless server does not remember anything from one request to another. For example, a HTTP server is stateless.
  - Stateful servers maintains information about its clients.

# Design Issues for Servers

## ➤ Where clients contact a server?

- In all cases, clients send requests to an end point, also called a port, at the machine where the server is running.
- Each server listens to a specific end point:
  - Servers that handle Internet FTP requests always listen to TCP **port 21**.
  - An HTTP server for the www listen to TCP port 80.

## ➤ How to handle communication interrupts?

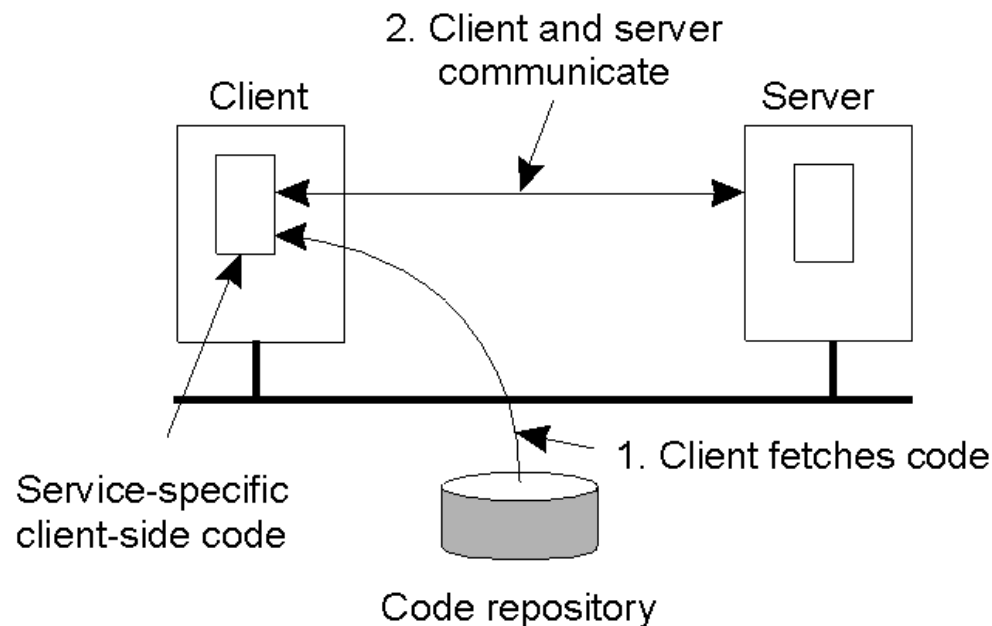
- Use **out-of-band** data. Example: to cancel the upload of a huge file.
- Server listens to separate endpoint, which has higher priority, while also listening to the normal endpoint (with lower priority).
- Send urgent data on the same connection. Can be done with TCP, where the server gets a signal (SIGURG) on receiving urgent data.

# Reasons for Code Migration

- Improve computing performance by moving processes from heavily-loaded machines to lightly loaded machines.
  - E.g. Java applets
- Improve communication times by shipping code to systems where large data sets reside.
  - E.g. a client ships code to a database server or vice versa.

# Reasons for Code Migration

- Flexibility to dynamically configure distributed systems.

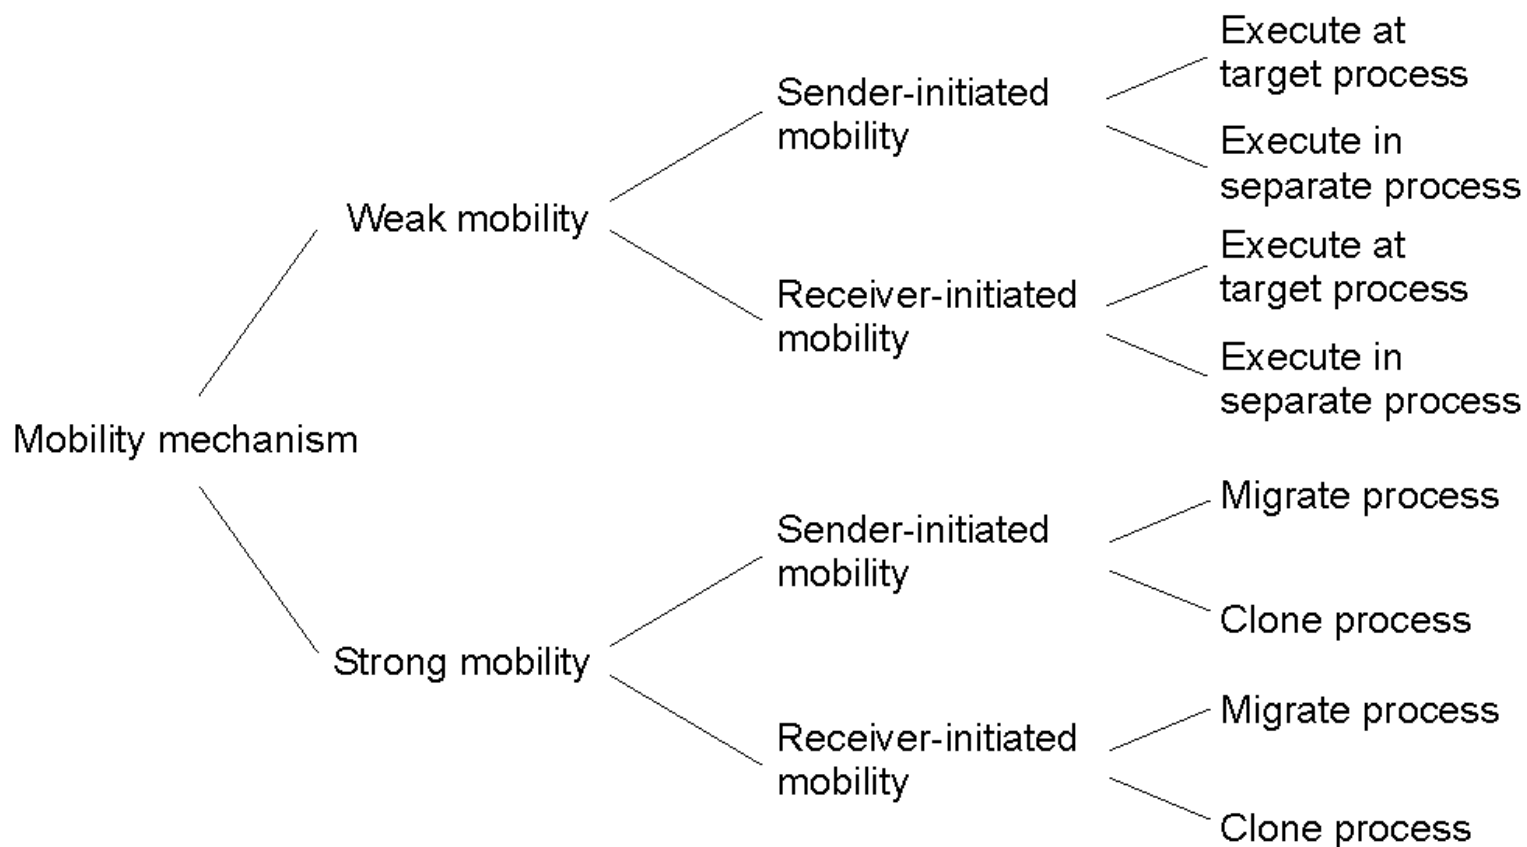


The principle of dynamically configuring a client to communicate to a server. The client first fetches the necessary software, and then invokes the server

# Migration Models

- A process consists of three segments: code segment, resource segment, execution segment.
- Migration (or mobility) can be **weak** or **strong**:
  - **Weak migration**: only the code segment can be transferred. Transferred program always starts at one of the predefined starting positions. E.g. java applets.
  - **Strong migration**: code and execution segments can both be transferred.
- Also, migration can be **sender-initiated** or **receiver-initiated**:
  - **Sender initiated**: uploading code to a server. Ex: Client sending a query to database server.
  - **Receiver-initiated**: downloading code from a server by a client. Ex: client downloading Java applets from a server.

# Migration Models



# Resource Migration

- Depends on type of “resource”:
  - By identifier: specific web site, ftp server
  - By value: Java libraries
  - By type: printers, local devices
- Depends on type of “attachments”
  - Unattached to any node: data files
  - Fastened resources (can be moved only at high cost)
    - Ex: database, web sites
  - Fixed resources
    - Ex: local devices, communication end points



# Resource Migration

Process-to-Resource Binding	Resource-To-Machine Binding		
	Unattached	Fastened	Fixed
By identifier	MV (or GR)	GR (or MV)	GR
By value	CP ( or MV, GR)	GR (or CP)	GR
By type	RB (or MV, CP)	RB (or GR, CP)	RB (or GR)

- Actions to be taken with respect to the references to local resources when migrating code to another machine
  - **GR:** Establish a global system-wide reference
  - **MV:** Move the resource
  - **CP:** Copy the value of the resource
  - **RB:** Rebind the process to a locally available resource

# Activity



# What happens to a TCP port opened by a migrating process?

- A. GR: Establish a global system-wide reference
- B. MV: Move the resource
- C. CP: Copy the value of the resource
- D. RB: Rebind the process to a locally available resource

# What happen to a reference to a File when the code is moved?

- A. GR: Establish a global system-wide reference
- B. MV: Move the resource
- C. CP: Copy the value of the resource
- D. RB: Rebind the process to a locally available resource

# Next Chapter

➤ Communication

**Questions?!**

# Please rate your understanding of the chapter

I understand most of the chapter topics (Threads, Virtualization, Server, Code Migration)

- A. Strongly Agree
- B. Agree
- C. Disagree
- D. Strongly Disagree