


Program – 3

AIM: Implement Mutual Exclusion using Centralized Algorithm.

Introduction and Theory

In centralized algorithm one process is elected as the coordinator which may be the machine. Whenever a process wants to enter a critical region, it sends a request message to the coordinator stating which critical region it wants to enter and asking for permission. If no other process is currently in that critical region, the coordinator sends back a reply granting permission. When the reply arrives, the requesting process enters the critical region. When another process asks for permission to enter the same critical region. Now the coordinator knows that a different process is already in the critical region, so it cannot grant permission. The coordinator just refrains from replying, thus blocking process 2, which is waiting for a reply or it could send a reply 'permission denied.' When process 1 exits the critical region, it sends a message to the coordinator releasing its exclusive access. The coordinator takes the first item off the queue of deferred requests and sends that process a grant message. If the process was still blocked it unblocks and enters the critical region. If an explicit message has already been sent denying permission, the process will have to poll for incoming traffic or block later. When it sees the grant, it can enter the critical region.



- Advantages
 - Algorithm guarantees mutual exclusion by letting one process at a time into each critical region.
 - It is also fair as requests are granted in the order in which they are received.
 - No process ever waits forever so no starvation.
 - Easy to implement so it requires only three messages per use of a critical region (request, grant, release).
 - Used for more general resource allocation rather than just managing critical regions.
- Disadvantages
 - The coordinator is a single point of failure, the entire system may go down if it crashes.
 - If processes normally block after making a request, they cannot distinguish a dead coordinator from "permission denied" since no message comes back.
 - In a large system a single coordinator can become a performance bottleneck.

Program –3

Code

Client

```
1  #include <sys/socket.h>
2  #include <sys/types.h>
3  #include <netinet/in.h>
4  #include <netdb.h>
5  #include <stdio.h>
6  #include <string.h>
7  #include <stdlib.h>
8  #include <unistd.h>
9  #include <errno.h>
10 #include <arpa/inet.h>
11 #include <unistd.h>
12 typedef struct resources
13 {
14     int A;
15     char B;
16     int C;
17     char D;
18 }resources;
19 int main()
20 {
21     struct sockaddr_in sa; // Socket address data structure
22     resources R;
23     int n, sockfd; // read and source
24     char buff[1025], obuff[256]; // buffer to store the read
25     stream
26     int snded, rec;
27
28     sockfd = socket(PF_INET, SOCK_STREAM, 0); // New socket
29     created
30     // Checking for valid socket
31     if (sockfd < 0)
32     {
33         printf("Error in creation\n");
34         exit(0);
35     }
36     else
37         printf("Socket created\n");
38
39     // Clearing and assigning type and address to the socket
40     bzero(&sa, sizeof(sa));
41     sa.sin_family = AF_INET;
42     sa.sin_port = htons(8888);
43
44     // establishing and verifying the connection
45     if (connect(sockfd, (struct sockaddr_in*)&sa, sizeof(sa)) <
46     0)
47     {
48         printf("Connection failed\n");
49         exit(0);
50     }
51     else
52         printf("Connection made\n");
53 }
```

Program – 3

```

54     while (1)
55     {
56         snded = write(sockfd, "PING", 5);
57         if (snded > -1)
58             printf("SENT PING\n");
59         rec = read(sockfd, obuff, 256);
60         obuff[rec] = '\0';
61         if (strcmp(obuff, "PONG") == 0)
62         {
63             usleep(750);
64             FILE *f;
65             f = fopen("shared_mem.txt", "r");
66             fread(&R, sizeof(R), 1, f);
67             fclose(f);
68             printf("read %d, %d, %d, %d from server\n", R.A, R.B,
69 R.C, R.D );
70             R.A += 1;
71             R.B += 1;
72             R.C += 1;
73             R.D += 1;
74             f = fopen("shared_mem.txt", "w");
75             fwrite(&R, sizeof(R), 1, f);
76             fclose(f);
77             printf("Got access to CS\n");
78             snded = write(sockfd, "DONE", 4);
79             printf("Freeing Lock\n");
80             break;
81         }
82     }
83     // Reading and priting data from the server after
84 verification
85     close(sockfd); // Closing the socket
86     return 0;

```

Controller

```

1  #include <sys/socket.h>
2  #include <netinet/in.h>
3  #include <arpa/inet.h>
4  #include <stdio.h>
5  #include <stdlib.h>
6  #include <unistd.h>
7  #include <errno.h>
8  #include <string.h>
9  #include <sys/types.h>
10 #include <time.h>
11 #define TRUE 1
12 #define FALSE 0
13 typedef struct resources{
14     int A;
15     int B;
16     int C;
17     int D;
18 }resources;
19 int main(){
20     resources R, temp;

```

Program – 3

```

21     R.A = 1;
22     R.B = 2;
23     R.C = 3;
24     R.D = 4;
25     FILE *file;
26     file = fopen("shared_mem.txt", "w");
27     fwrite(&R, sizeof(R), 1, file);
28     fclose(file);
29     struct sockaddr_in sa; // Socket address data structure
30     int opt = TRUE, addrlen;
31     int sockfd, clients[50]; // Source and destination addresses
32     char buff[256]; // Buffer to hold the out-going stream
33     int rec, i, sd, activity, new_sock, sendd;
34     int max_sd;
35     int flag = 0;
36     sockfd = socket(AF_INET, SOCK_STREAM, 0); // New socket
37     created
38
39         // Checking for valid socket
40     memset(clients, 0, sizeof(clients));
41
42     fd_set readfds;
43     if (sockfd < 0)
44     {
45         printf("Error in creating socket\n");
46         exit(0);
47     }
48     else
49     {
50         printf("Socket Created\n");
51     }
52     if (setsockopt(sockfd, SOL_SOCKET, SO_REUSEADDR, (char
53 *)&opt, sizeof(opt)) < 0)
54     {
55         printf("error\n");
56     }
57     // Clearing and assigning type and address to the socket
58     printf("Socket created\n");
59     bzero(&sa, sizeof(sa));
60     sa.sin_family = AF_INET;
61     sa.sin_port = htons(8888);
62     sa.sin_addr.s_addr = htonl(INADDR_ANY);
63
64     // binding and verifying the socket to address
65     if (bind(sockfd, (struct sockaddr*)&sa, sizeof(sa))<0)
66     {
67         printf("Bind Error\n");
68     }
69     else
70         printf("Binded\n");
71     // starts the server with a max client queue size set as 10
72     listen(sockfd, 10);
73     addrlen = sizeof(sa);
74     // server run
75     while (TRUE)
76     {
77         // Clearing socket set

```

Program – 3

```

78         FD_ZERO(&readfds);
79
80         FD_SET(sockfd, &readfds);
81         max_sd = sockfd;
82         for (i = 0; i < 50; i++){
83             sd = clients[i];
84             if (sd > 0)
85                 FD_SET(sd, &readfds);
86             if (sd > max_sd)
87                 max_sd = sd;
88         }
89         activity = select(max_sd + 1, &readfds, NULL, NULL,
90 NULL);
91         if (activity < 0)
92             printf("Select error\n");
93         if (FD_ISSET(sockfd, &readfds))
94         {
95             if ((new_sock = accept(sockfd, (struct
96 sockaddr*)&NULL, NULL)) < 0)
97                 perror("accept");
98             else
99             {
100                 printf("New connection, sock fd
101 %d\n", new_sock);
102             }
103             sended = send(new_sock, buff, strlen(buff),
104 0);
105             if (sended < 0)
106                 perror("Send");
107             for (i = 0; i < 50; i++){
108                 if (clients[i] == 0){
109                     clients[i] = new_sock;
110                     break;
111                 }
112             }
113         }
114         for (i = 0; i < 50; i++){
115             sd = clients[i];
116             if (FD_ISSET(sd, &readfds))
117             {
118                 FILE *file;
119                 file = fopen("shared_mem.txt", "r");
120                 fread(&temp, sizeof(temp), 1, file);
121                 fclose(file);
122                 rec = read(sd, buff, 256);
123                 if (rec == 0){
124                     getpeername(sd, (struct
125 sockaddr*)&sa, \
126                                     (socklen_t*)&sa);
127                     printf("%d has disconnected
128 unexpectedly with ip %s and port %d\n", sd, inet_ntoa(sa.sin_addr),
129 ntohs(sa.sin_port));
130                     printf("recovering data\n");
131                     FILE *file;
132                     file = fopen("shared_mem.txt", "w+");
133                     fwrite(&temp, sizeof(temp), 1, file);
134                     fclose(file);

```

Program – 3

```

135         close(sd);
136         clients[i] = 0;
137     }
138     else
139     {
140         buff[rec] = '\0';
141         printf("recieved %s from
142 %d\n", buff, sd);
143
144         if (strcmp(buff, "PING") == 0 && flag == 1)
145         {
146             printf("Read buffer = %s, from %d and send
147 NACK\n", buff, sd);
148             send = write(sd, "NACK", 4);
149         }
150
151         else if (strcmp(buff, "PING") == 0 && flag ==
152 0) {
153             printf("Read Buffer = %s, from %d\n", buff,
154 sd);
155             flag = 1;
156             send = write(sd, "PONG", 4);
157             // rec = read(sd, buff, 4);
158             // if (rec > 0 && strcmp(buff, "DONE") == 0)
159             // {
160             //     printf("Lock freed\n");
161             //     flag = 0;
162             //     FILE *f1;
163             //     f1 = fopen("shared_mem.txt", "r");
164             //     fread(&temp, sizeof(temp), 1, f1);
165             //     printf("Read %d, %d, %d, %d from
166 %d\n", temp.A, temp.B, temp.C, temp.D, sd );
167             //     fclose(f1);
168             //     clients[i] = 0;
169             //     close(sd);
170             //     break;
171             // }
172         }
173         else if (strcmp(buff, "DONE") == 0) {
174             printf("Lock freed\n");
175             flag = 0;
176             FILE *f1;
177             f1 = fopen("shared_mem.txt", "r");
178             fread(&temp, sizeof(temp), 1, f1);
179             printf("Read %d, %d, %d, %d from %d\n",
180 temp.A, temp.B, temp.C, temp.D, sd );
181             fclose(f1);
182             clients[i] = 0;
183             close(sd);
184             break;
185         }
186     }
187 }
188 }
189 }
190 }
191 close(sockfd); // close the socket

```

Program – 3

```
192         return 0;
193     }
```

Results and Outputs:

```
c.Semesters/College.Stuff.Academic.Semesters.YEAR_4/SEM 7/C0403_Distributed_Systems/DisLAB$ ./outs/l2server
Socket Created
Socket created
Binded
New connection, sock fd 4
recieved PING from 4
Read Buffer = PING, from 4
recieved DONE from 4
Lock freed
Read 2, 3, 4, 5 from 4
New connection, sock fd 4
recieved PING from 4
Read Buffer = PING, from 4
recieved DONE from 4
Lock freed
Read 3, 4, 5, 6 from 4
New connection, sock fd 4
recieved PING from 4
Read Buffer = PING, from 4
recieved DONE from 4
Lock freed
Read 4, 5, 6, 7 from 4
```

Figure 1 Controller

Figure 2 Clients

```
rinzler@Jarvis:/mnt/h/College stuff/College Stuff.Academic/College Stuff.Academic.Semesters/College.Stuff.Academic.Semesters.YEAR_4/SEM 7/C0403_Distributed_Systems/DisLAB$ ./outs/l2client
Socket created
Connection made
SENT PING
read 1, 2, 3, 4 from server
Got access to CS
Freeing Lock
```

```
rinzler@Jarvis:/mnt/h/College stuff/College Stuff.Academic/College Stuff.Academic.Semesters/College.Stuff.Academic.Semesters.YEAR_4/SEM 7/C0403_Distributed_Systems/DisLAB$ ./outs/l2client
Socket created
Connection made
SENT PING
read 1, 2, 3, 4 from server
Got access to CS
Freeing Lock
```

```
rinzler@Jarvis:/mnt/h/College stuff/College Stuff.Academic/College Stuff.Academic.Semesters/College.Stuff.Academic.Semesters.YEAR_4/SEM 7/C0403_Distributed_Systems/DisLAB$ ./outs/l2client
Socket created
Connection made
SENT PING
read 1, 2, 3, 4 from server
Got access to CS
Freeing Lock
```

Program – 3

Findings and Learnings:

1. We successfully implemented Centralized Mutual Exclusion.
2. The system is better at handling crashed processes
3. The Entire system fails for Crashed Clients.

