

Distributed Systems

Chapter 3 Distributed Objects and Remote Invocation

5.1 Communication between distributed objects

5.2 Remote procedure call

5.3 Events and notifications

5.4 Case study: Java RMI

5.5 Summary

5.1 Communication between Distributed Objects

Topics to be addressed

- Distributed objects
- The distributed object model
- Design issues
- Implementation
- Distributed garbage collection

5.2.2 Distributed Objects

- Physical distribution of objects into different processes or computers in a distributed system
 - Object state consists of the values of its instance variables
 - Object methods invoked by remote method invocation (RMI)
 - Object encapsulation: object state accessed only by the object methods
- Usually adopt the client-server architecture
 - Basic model
- Objects are managed by servers and
- Their clients invoke their methods using RMI
- Steps
 1. The client sends the RMI request in a message to the server
 2. The server executes the invoked method of the object
 3. The server returns the result to the client in another message
- Other models

- Chains of related invocations: objects in servers may become clients of objects in other servers
- Object replication: objects can be replicated for fault tolerance and performance
- Object migration: objects can be migrated to enhancing performance and availability

5.2.3 The Distributed Object Model

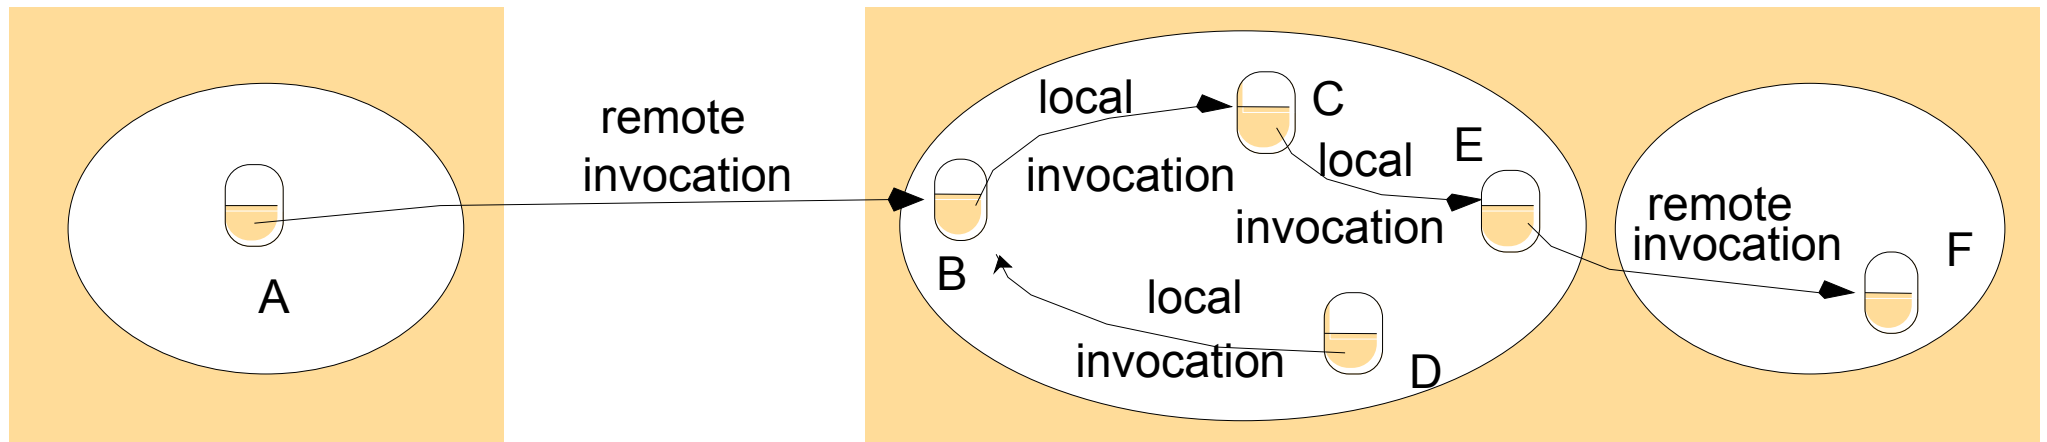


Figure 5.3 Remote and local method invocations

Two fundamental concepts: Remote Object Reference and Remote Interface

- Each process contains objects, some of which can receive remote invocations are called remote objects (B, F), others only local invocations
- Objects need to know the remote object reference of an object in another process in order to invoke its methods, called remote method invocations

- Every remote object has a remote interface that specifies which of its methods can be invoked remotely

Five Parts of Distributed Object Model

- Remote Object References

- accessing the remote object
- identifier throughout a distributed system
- can be passed as arguments

- Remote Interfaces

- specifying which methods can be invoked remotely
- name, arguments, return type
- Interface Definition Language (IDL) used for defining remote interface

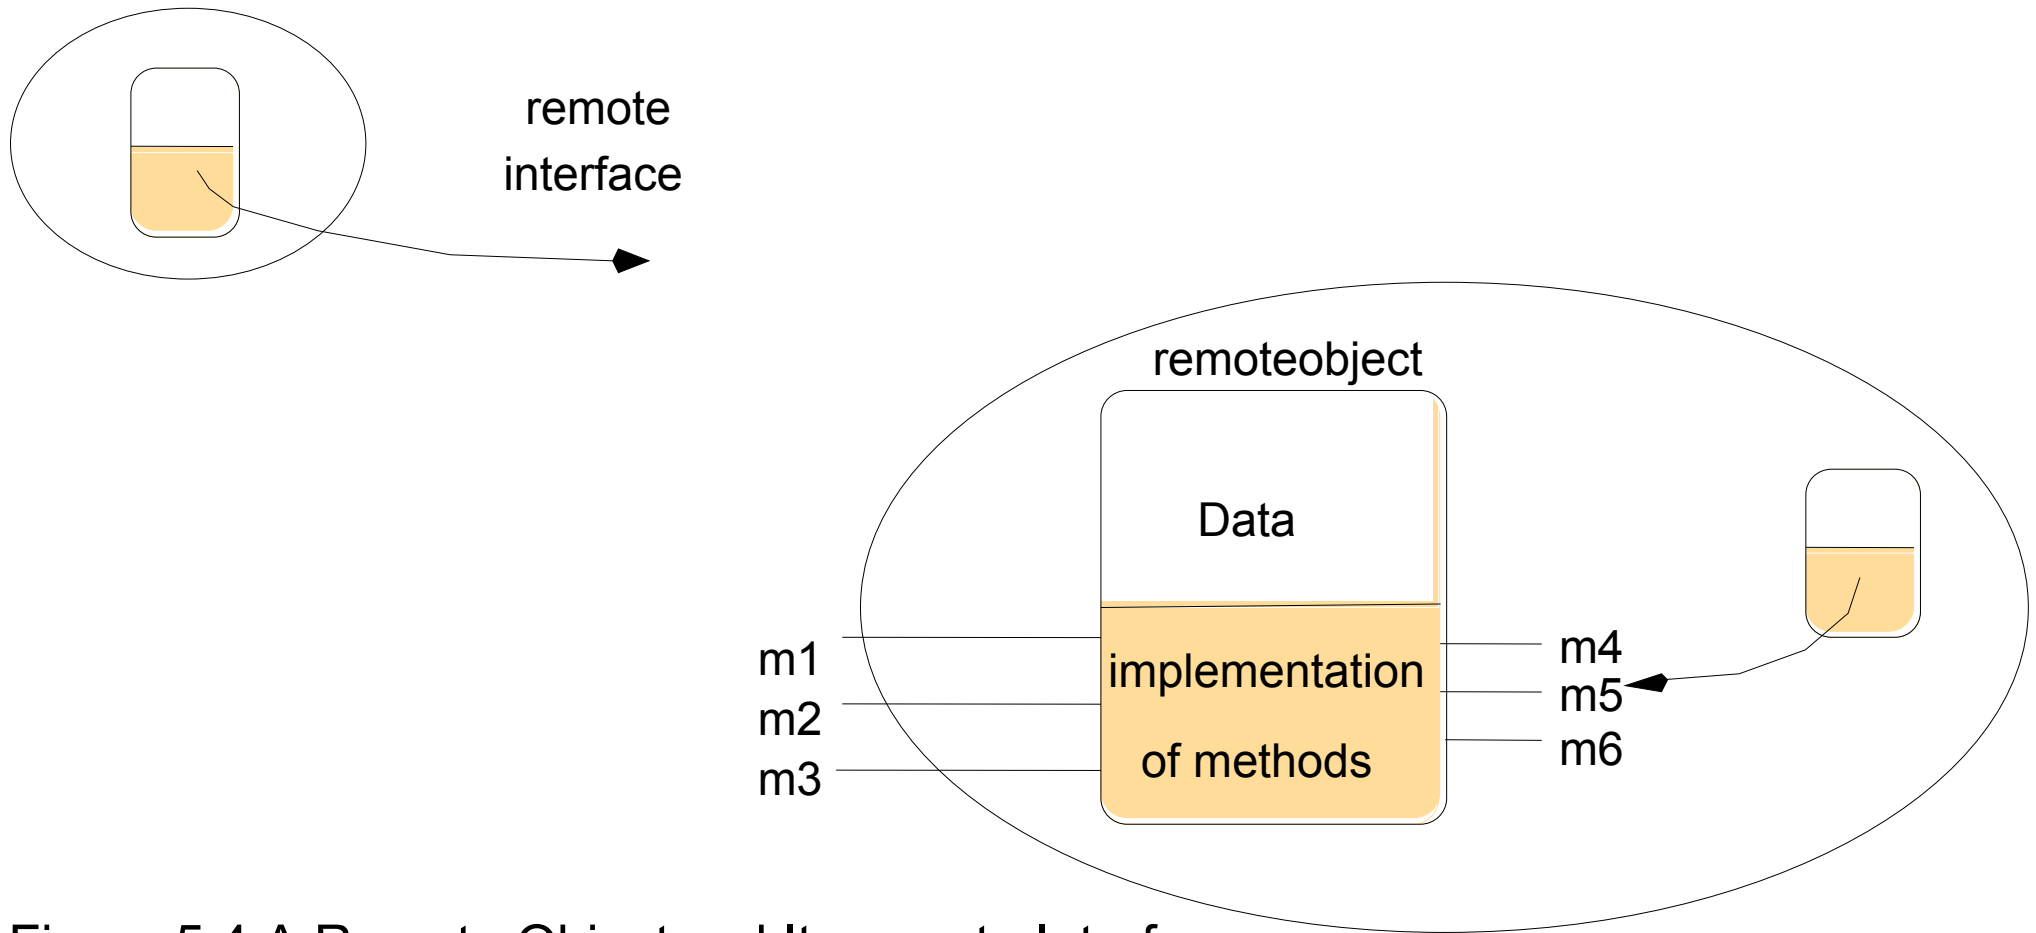


Figure 5.4 A Remote Object and Its remote Interface

{

Five Parts of Distributed Object Model (cont.)

- Actions

- An action initiated by a method invocation may result in further invocations on methods in other objects located in different processes or computers
- Remote invocations could lead to the instantiation of new objects, ie. objects M and N of Figure 5.5

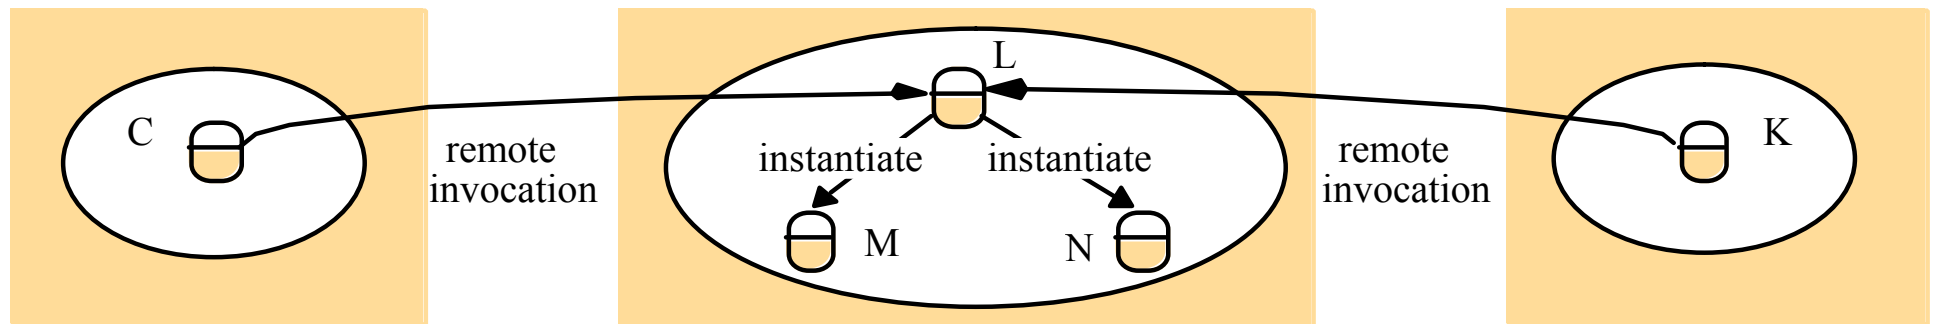
- Exceptions

- More kinds of exceptions: i.e. timeout exception

- RMI should be able to raise exceptions such as timeouts that are due to distribution as well as those raised during the execution of the method invoked

- Garbage Collection

- Distributed garbage collection is generally achieved by cooperation between the existing local garbage collector and an added module that carries out a form of distributed garbage collection, usually based on reference counting



5.2.4 Design Issues for RMI

- Two design issues that arise in extension of local method invocation for RMI
 - The choice of invocation semantics
 - Although local invocations are executed exactly once, this cannot always be the case for RMI due to transmission error
 - Either request or reply message may be lost
 - Either server or client may be crashed
 - The level of transparency
 - Make remote invocation as much like local invocation as possible

RMI Design Issues: Invocation Semantics

- Error handling for delivery guarantees
 - Retry request message: whether to retransmit the request message until either a reply is received or the server is assumed to have failed
 - Duplicate filtering: when retransmissions are used, whether to filter out duplicate requests at the server
 - Retransmission of results: whether to keep a history of result messages to enable lost results to be retransmitted without re-executing the operations
- Choices of invocation semantics
 - Maybe: the method executed once or not at all (no retry nor retransmit)
 - At-least-once: the method executed **at least** once
 - At-most-once: the method executed **exactly** once

Fault tolerance measures Invocation

*Retransmit request Duplicate Re-execute procedure semantics message
filtering or retransmit reply*

No Not applicable Not applicable *Maybe*

Yes No Re-execute procedure *At-least-once*

Yes Yes Retransmit reply *At-most-once*

Figure 5.6 Invocation semantics: choices of interest

RMI Design Issues: Transparency

- Transparent remote invocation: like a local call
 - marshalling/unmarshalling
 - locating remote objects
 - accessing/syntax
- Differences between local and remote invocations
 - latency: a remote invocation is usually several order of magnitude greater than that of a local one

- availability: remote invocation is more likely to fail
- errors/exceptions: failure of the network? server? hard to tell
- syntax might need to be different to handle different local vs remote errors/exceptions (e.g. Argus)
 - consistency on the remote machine:
- Argus: incomplete transactions, abort, restore states [as if the call was never made]

5.2.5 Implementation of RMI

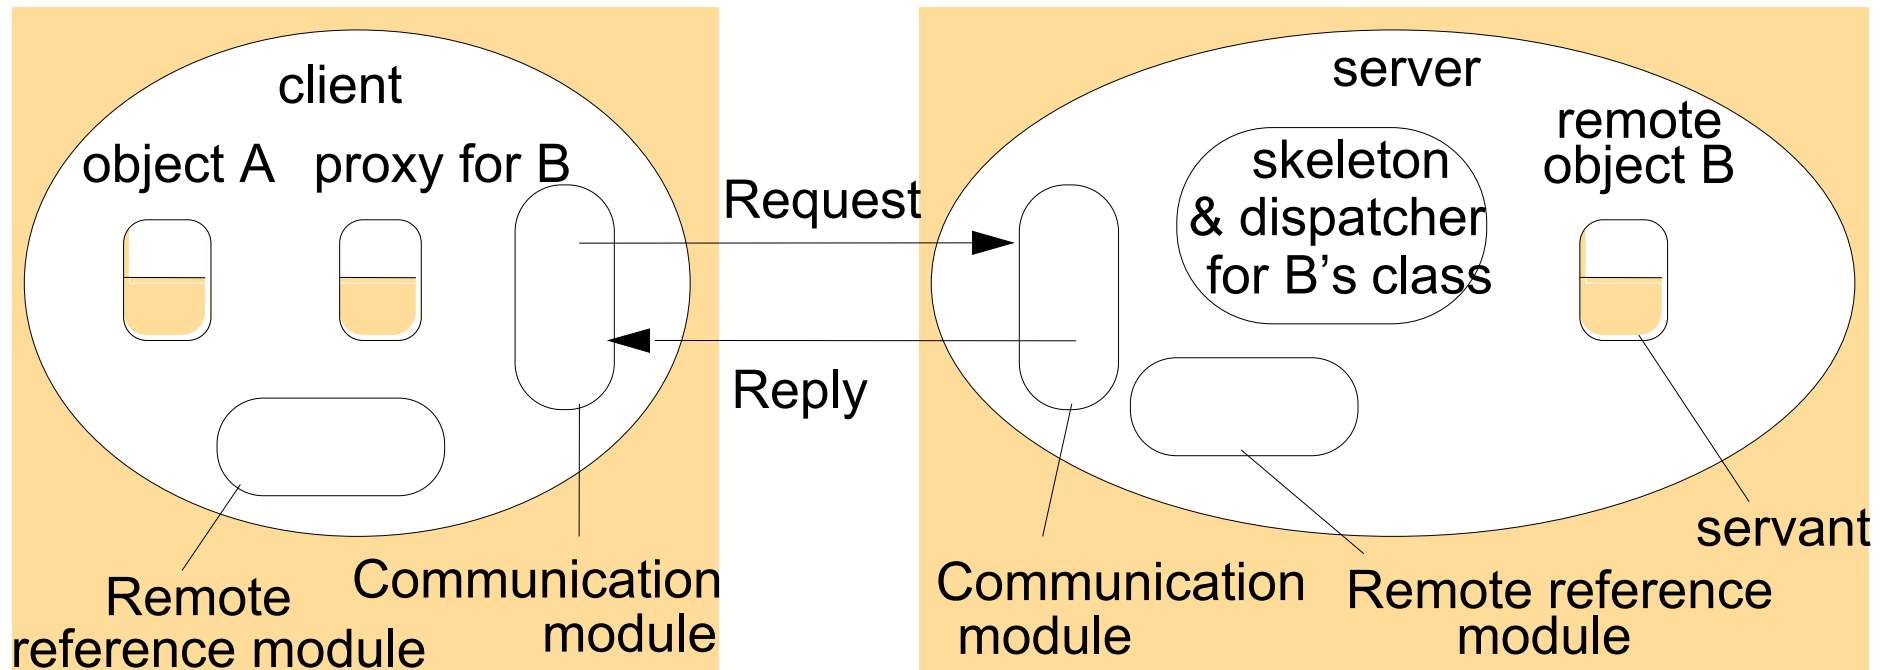


Figure 5.7 Role of proxy and skeleton in RMI

•Communication module

- Two cooperating communication modules carry out the request-reply protocols: message type, requestID, remote object reference
 - Transmit request and reply messages between client and server
 - Implement specific invocation semantics

–The communication module in the server

- selects the dispatcher for the class of the object to be invoked,
- passes on local reference from remote reference module, • returns request

Implementation of RMI (cont.)

- Remote reference module
 - Responsible for translating between local and remote object references and for creating remote object references
 - remote object table: records the correspondence between local and remote object references
 - remote objects held by the process (B on server)
 - local proxy (B on client)
 - When a remote object is to be passed for the first time, the module is asked to create a remote object reference, which it adds to its table
- Servant
 - An instance of a class which provides the body of a remote object – handles the remote requests

- RMI software
 - Proxy: behaves like a local object, but represents the remote object
 - Dispatcher: look at the methodID and call the corresponding method in the skeleton
 - Skeleton: implements the method
- Generated automatically by an interface compiler

Implementation Alternatives of RMI

- Dynamic invocation
 - Proxies are static—interface compiled into client code
 - Dynamic—interface available during run time
- Generic invocation; more info in “Interface Repository” (COBRA)
- Dynamic loading of classes (Java RMI)
- Binder
 - A separate service to locate service/object by name through table mapping for names and remote object references
- Activation of remote objects

- Motivation: many server objects not necessarily in use all of the time
 - Servers can be started whenever they are needed by clients, similar to inetd–
- Object status: active or passive
 - active: available for invocation in a running process
 - passive: not running, state is stored and methods are pending – Activation of objects:
 - creating an active object from the corresponding passive object by creating a new instance of its class
 - initializing its instance variables from the stored state
- Responsibilities of *activator*
 - Register passive objects that are available for activation
 - Start named server processes and activate remote objects in them
 - Keep track of the locations of the servers for remote objects that it has already activated

Implementation Alternatives of RMI (cont.)

- Persistent object stores
 - An object that is guaranteed to live between activations of processes is called a *persistent object*
 - Persistent object store: managing the persistent objects
- stored in marshaled form on disk for retrieval

- saved those that were modified
- Deciding whether an object is persistent or not:
 - persistent root: any descendent objects are persistent (persistent Java, PerDiS)
 - some classes are declared persistent (Arjuna system)
- Object location
- specifying a location: ip address, port #, ...
- location service for migratable objects
 - Map remote object references to their probable current locations
 - Cache/broadcast scheme (similar to ARP)
- Cache locations
- If not in cache, broadcast to find it
 - Improvement: forwarding (similar to mobile IP)

5.2.6 Distributed Garbage Collection

- Aim: ensure that an object
 - continues to exist if a local or remote reference to it is still held anywhere
 - be collected as soon as no object any longer holds a reference to it

- General approach: reference count
- Java's approach
 - the server of an object (B) keeps track of proxies
 - when a proxy is created for a remote object
- addRef(B) tells the server to add an entry
 - when the local host's garbage collector removes the proxy
 - removeRef(B) tells the server to remove the entry
 - when no entries for object B, the object on server is deallocated

5.3 Remote Procedure Call

- client:
"stub"
instead of
"proxy"

(same
function,
different
names) –
local call,
marshal
arguments,
communicat
e the
request

- **server:** –
dispatcher

- "stub": unmarshal arguments, communicate the results back

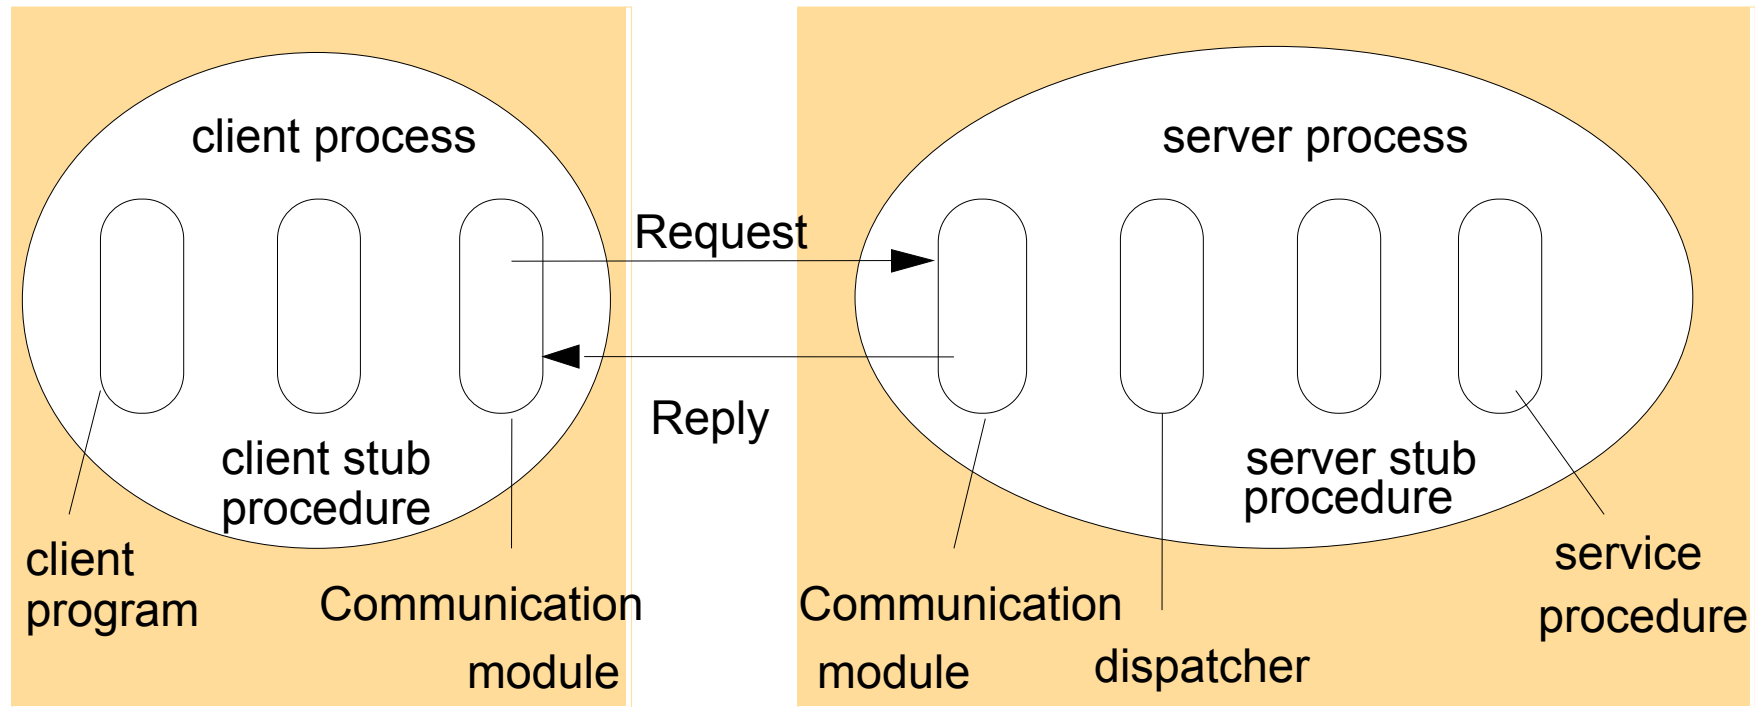


Figure 5.8 Role of client and server stub procedures in RPC in the context of a procedural language

Distributed Event Notification

- Distributed event notification – decouple publishers from subscribers via an event service (manager)
- Architecture: roles of participating objects
 - object of interest (usually changes in states are interesting)
 - event
 - notification
 - subscriber
 - observer object (proxy) [reduce work on the object of interest]
- forwarding
- filtering of events types and content/attributes
- patterns of events (occurrence of multiple events, not just one)
- mailboxes (notifications in batches, subscriber might not be ready)
- publisher (object of interest or observer object)
 - generates event notifications

Example: Distributed Event Notification

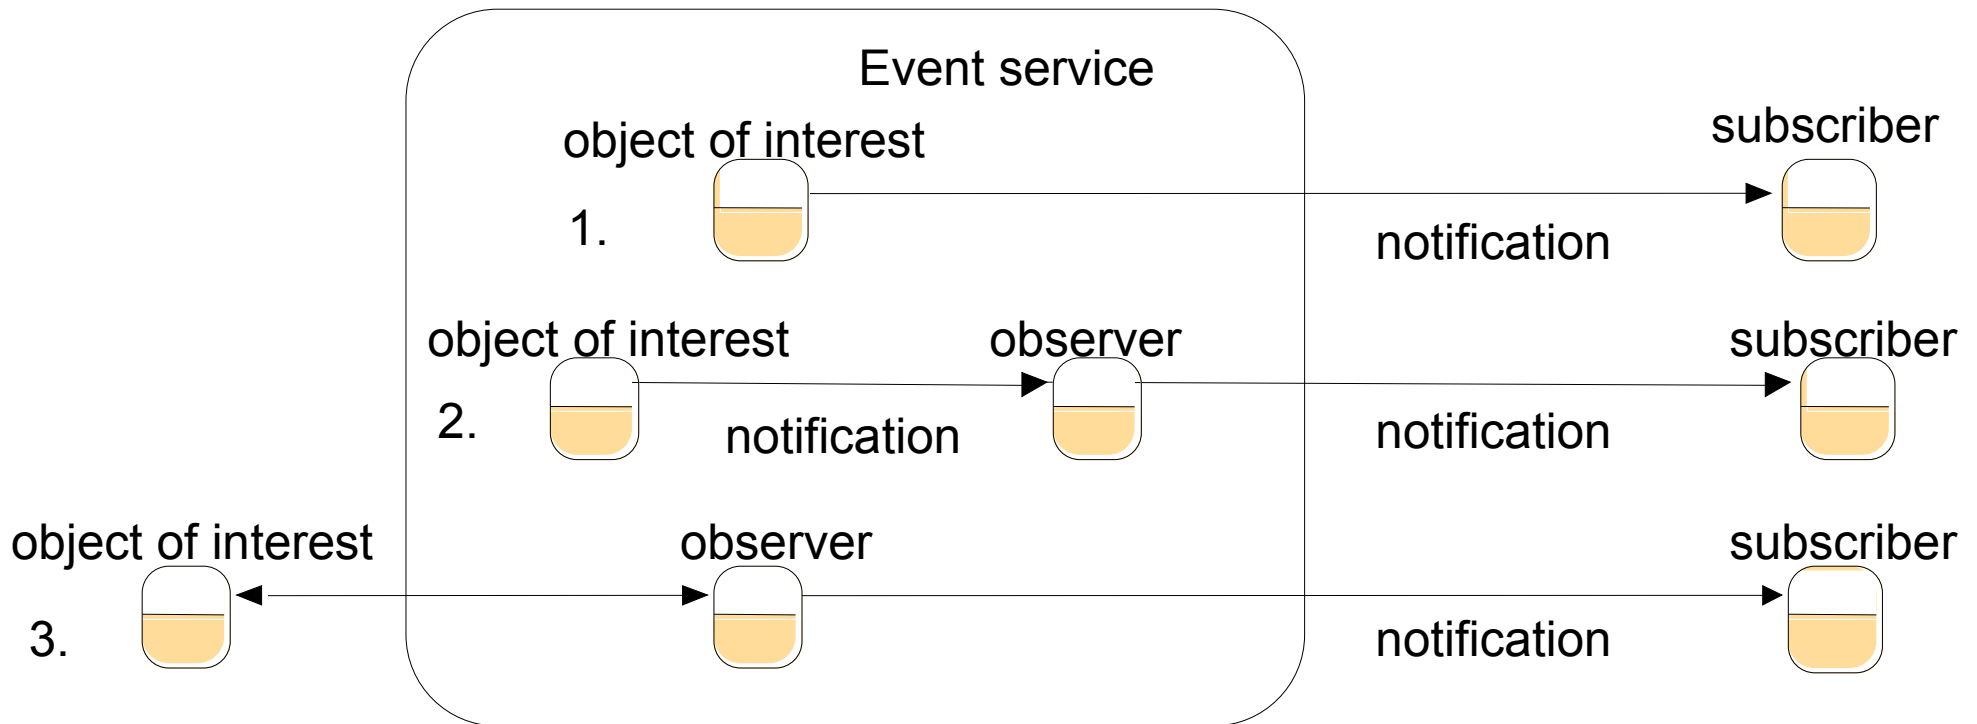


Figure 5.11 Architecture for distributed event notification

- Three cases
 - Inside object without an observer: send notifications directly to the subscribers

- Inside object with an observer: send notification via the observer to the subscribers
 - Outside object (with an observer)
 1. an observer queries the object of interest in order to discover when events occur
 2. The observer sends notifications to the subscribers